



STEVENS INSTITUTE OF TECHNOLOGY



# Natural Language Processing (NLP)

Rensheng Wang,  
<https://sit.instructure.com/courses/29876>  
January 24, 2019



# What is text?

You can think of text as a sequence of

- ☐ Characters
- ☐ Words
- ☐ Phrases and named entities
- ☐ Sentences
- ☐ Paragraphs
- ☐ ...



# What is a word?

↳ a meaningful structured unit

❑ It seems natural to think of a text as a sequence of words

👉 A word is a meaningful sequence of characters

❑ How to find the boundaries of words?

👉 In English we can split a sentence by spaces or punctuation

❑ **Input:**

Fridends, No smoking, Let us talk;

❑ **Output:**

Friends No smoking Let us talk



# Tokenization

Tokenization is a process that splits an input sequence into so-called tokens

- ❑ You can think of a token as a useful unit for semantic processing
- ❑ Can be a word, sentence, paragraph, etc.



# Tokenization

Tokenization is a process that splits an input sequence into so-called tokens

- You can think of a token as a useful unit for semantic processing
- Can be a word, sentence, paragraph, etc.

An example of simple whitespace tokenizer

□ `nltk.tokenize.WhitespaceTokenizer`


This is Tom's book, isn't it?

👉 Problem: “it” and “it?” are different tokens with the same meaning




# Tokenization

Lets try to also split by punctuation

 `nlTK.tokenize.WordPunctTokenizer`

This is Tom 's book , isn ' t it ?

 Problem: “s”, “isn”, “t” are not very meaningful



# Tokenization

Lets try to also split by punctuation

```
nltk.tokenize.WordPunctTokenizer
```

This is Tom 's book , isn ' t it ?

👉 Problem: “s”, “isn”, “t” are not very meaningful

We can come up with a set of rules

```
nltk.tokenize.TreebankWordTokenizer
```

This is Tom 's book , is n't it ?

👉 “s” and “n't” are more meaningful for processing



# Python tokenization example

```
□ import nltk
```





## Token normalization

We may want the same token for different forms of the word

❑ wolf, wolves → wolf

❑ talk, talks → talk

### Stemming

❑ A process of removing and replacing suffixes to get to the root form of the word, which is called the **stem**

❑ Usually refers to heuristics that chop off suffixes

### Lemmatization

❑ Usually refers to doing things properly with the use of a vocabulary and morphological analysis

❑ Returns the base or dictionary form of a word, which is known as the **lemma**



# Stemming example

## Porters stemmer

□ 5 heuristic phases of word reductions, applied sequentially

□ Example of phase 1 rules:

Rule	Example
SSSES → SS	caresses → caress
IES → I	ponies → poni
SS → SS	caress → caress
S →	cats → cat

□ `nlTK.stem.PorterStemmer`

□ Examples:

feet → feet	cats → cat
wolves → wolv	talked → talk

□ Problem: fails on irregular forms, produces non-words



# Lemmatization example

## WordNet lemmatizer

- Uses the WordNet Database to lookup lemmas

- `nltk.stem.WordNetLemmatizer`

- Examples:

feet → foot	cats → cat
wolves → wolf	talked → talked

- Problems: not all forms are reduced



# Lemmatization example

## WordNet lemmatizer

- Uses the WordNet Database to lookup lemmas

- `nltk.stem.WordNetLemmatizer`

- Examples:

feet → foot	cats → cat
wolves → wolf	talked → talked

- Problems: not all forms are reduced

👉 Conclusion: we need to try stemming or lemmatization and choose best for our task



# Python stemming example



## Further normalization

### Normalizing capital letters

- ☐ Us, us  $\rightarrow$  us (if both are pronoun)
- ☐ us, US (could be pronoun and country)
- ☐ We can use heuristics:
  - ☐ lowercasing the beginning of the sentence
  - ☐ lowercasing words in titles
  - ☐ leave mid-sentence words as they are
- ☐ Or we can use machine learning to retrieve true casing  $\rightarrow$  hard

### Acronyms

- ☐ eta, e.t.a., E.T.A.  $\rightarrow$  E.T.A.
- ☐ We can write a bunch of regular expressions  $\rightarrow$  hard



# Transforming tokens into features: BOW

## Bag of words (BOW)

- Lets count occurrences of a particular token in our text
- Motivation: we are looking for marker words like “excellent” or “disappointed”
- For each token we will have a feature column, this is called text vectorization.

good movie
not a good movie
did not like

⇒

good	movie	not	a	did	like
1	1	0	0	0	0
1	1	1	1	0	0
0	0	1	0	1	1

- Problems:

- we loose word order, hence the name “bag of words”
- counters are not normalized

if it occurs  
0 if not



## Lets preserve some ordering

We can count token pairs, triplets, etc.

*order of the sentence*

□ Also known as **n-grams**

- 1-grams for tokens
- 2-grams for token pairs
- ...

good movie
not a good movie
did not like

⇒

good movie	movie	did not	a	...
1	1	0	0	...
1	1	1	1	...
0	0	1	0	...

□ Problems:

- too many features





# Remove some n-grams

Lets remove some n-grams from features based on their occurrence frequency in documents of our corpus



## Remove some n-grams

Lets remove some n-grams from features based on their occurrence frequency in documents of our **corpus**

- ☐ High frequency n-grams:
  - Articles, prepositions, etc. (example: and, a, the)
  - They are called **stop-words**, they won't help us to discriminate texts → remove them
- ☐ Lowfrequencyn-grams:
  - Typos, rare n-grams
  - We dont need them either, otherwise we will likely **overfit**
- ☐ Mediumfrequencyn-grams:
  - Those are good n-grams



# There are a lot of medium frequency n-grams

- It proved to be useful to look at n-gram frequency in our corpus for filtering out bad n-grams
- What if we use it for ranking of medium frequency n-grams?
- **Idea:** the n-gram with smaller frequency can be more discriminating because it can capture a specific issue in the review



# TF-IDF

## Term frequency (TF)

□  $\text{tf}(t, d)$  – frequency for term (or n-gram)  $t$  in document  $d$

□ Variants:

weighting scheme	TF weight
binary	0, 1
raw count	$f_{t,d}$
term frequency	$f_{t,d} / \sum_{t' \in d} f_{t',d}$
log normalization	$1 + \log(f_{t,d})$



# TF-IDF

eg: finding no. of citations for a research paper

Inverse document frequency (IDF)

- $N = |D|$  : total number of documents in corpus
- $|\{d \in D : t \in d\}|$  : number of documents where the term  $t$  appears
- $\text{idf}(t, D) = \log \frac{N}{|\{d \in D : t \in d\}|}$

TF-IDF

- $\text{tfidf}(t, d, D) = \text{tf}(t, d) \cdot \text{idf}(t, D)$
- A high weight in TF-IDF is reached by a high term frequency (in the given document) and a low document frequency of the term in the whole collection of documents

local → global  
how freq. it appears?  
where?



# Better BOW

- ❑ Replace counters with TF-IDF
- ❑ Normalize the result row-wise (divide by  $L_2$ -norm)

good movie	⇒	good movie	movie	did not	...
not a good movie		0.17	0.17	0	...
did not like		0.17	0.17	0	...
		0	0	0.47	...



# Python TF-IDF example



# First text classification model: Sentiment classification

## IMDB movie reviews dataset

- ☐ <http://ai.stanford.edu/~amaas/data/sentiment/>
- ☐ Contains 25000 positive and 25000 negative reviews
- ☐ Contains at most 30 reviews per movie
- ☐ At least 7 stars out of 10  $\rightarrow$  positive (label = 1)
- ☐ At most 4 stars out of 10  $\rightarrow$  negative (label = 0)
- ☐ 50/50 train/test split
- ☐ Evaluation: accuracy





# Sentiment classification

Features: bag of 1-grams with TF-IDF values

- 25000 rows, 74849 columns for training
- Extremely sparse feature matrix : 99.8% are zeros

acting	actingjob	actings	actingwise
0.000000	0.0	0.0	0.0
0.000000	0.0	0.0	0.0
0.053504	0.0	0.0	0.0
0.033293	0.0	0.0	0.0
0.000000	0.0	0.0	0.0



# Sentiment classification

Model: Logistic regression

- ❑  $p(\mathbf{y} = 1|\mathbf{x}) = \sigma(\mathbf{w}^T \mathbf{x})$
- ❑ Linear classification model
- ❑ Can handle sparse data
- ❑ Fast to train
- ❑ Weights can be interpreted



# Sentiment classification

Logistic regression over bag of 1-grams with TF-IDF

□ Accuracy on test set: 88.5%

□ Let's look at learnt weights:

ngram	weight	VS	ngram	weight
great	9.042803		worst	-12.748257
excellent	8.487379		awful	-9.150810
perfect	6.907277		bad	-8.974974
best	6.440972		waste	-8.944854
wonderful	6.237365		boring	-8.340877
Top positive			Top negative	



## Better sentiment classification

Let us try to add 2-grams

- ❑ Throw away n-grams seen less than 5 times
- ❑ 25000 rows, 156821 columns for training

<b>and am</b>	<b>and amanda</b>	<b>and amateur</b>	<b>and amateurish</b>	<b>and amazing</b>
0.068255	0.0	0.0	0.0	0.0
0.000000	0.0	0.0	0.0	0.0
0.000000	0.0	0.0	0.0	0.0
0.000000	0.0	0.0	0.0	0.0
0.000000	0.0	0.0	0.0	0.0



## Better sentiment classification

Logistic regression over bag of 1,2-grams with TF-IDF

▣ Accuracy on test set: 89.9% (+1.5%)

▣ Lets look at learnt weights:

well worth 13.788515

best 13.633200

rare 13.570259

better than 13.500025

Near top positive

VS

bad -24.467648

poor -24.319746

the worst -23.773352

waste -22.880340

Near top negative



# How to make it even better

## □ Play around with tokenization

👉 Special tokens like emoji, :) and !!! can help Try to normalize tokens

## □ Adding stemming or lemmatization Try different models

👉 SVM, Naive Bayes, ...

## □ Throw BOW away and use Deep Learning

👉 <https://arxiv.org/pdf/1512.08183.pdf>

👉 Accuracy on test set in 2016: 92.14% (+2.5%)

